

Experiment - 5

Student Name: Abhi Kumar

UID: 23BCS12907

Branch: BE-CSE

Section/Group: KRG_2B

Semester: 5th

Date of Performance: 22/9/25

Subject Name: Advanced Database and Management System

Subject Code: 23CSP-333

1. Problem Description/Aim:

Medium-Problem Title: Generate 1 million records per ID in 'transaction_data' using generate_series() and random() ,create a normal view and a materialized view 'sales_summary' with aggregated metrics (total_quantity_sold , total_sales, total_orders) , and compare their performance and execution time.

Procedure (Step-by-Step):

1. Create a large dataset:
 - Create a table names transaction_data (id , value) with 1 million records.
 - take id 1 and 2, and for each id, generate 1 million records in value column
 - Use Generate_series () and random() to populate the data.
2. Create a normal view and materialized view to for sales_summary, which includes total_quantity_sold, total_sales, and total_orders with aggregation.
3. Compare the performance and execution time of both.

Sample Output Description:

The transaction_data table has 2 million rows (1 million per ID) with random values. The normal view sales_summary computes aggregates on the fly, while the materialized view sales_summary_mv stores precomputed results. Queries on the materialized view are much faster, but it needs refreshing when data changes, whereas the normal view always shows up-to-date results.

Hard-Problem Title: Create restricted views in the sales database to provide summarized, non-sensitive data to the reporting team, and control access using DCL commands(GRANT and REVOKE).

Procedure (Step-by-Step):

1. Create restricted views-
 - Define views that show only **aggregated sales data** (e.g., total_sales, total_orders) without exposing sensitive columns like customer details or payment info.

2. Assign access to reporting team(or client)-
 - Use “GRANT SELECT ON view_name TO reporting_user; “ to give access.
3. Revoke access if needed.
 - Use “REVOKE SELECT ON view_name FROM reporting_user;” to remove access.
4. Verify access
 - Reporting users can query the view but cannot access base tables directly, ensuring security.

Sample Output Description:

The result shows the restricted view providing summarized sales data only like

- Columns shown are - product_id,total_quantity_sold, total_sales, total_orders
- Columns hidden are - Customer names, addresses, payment details

A reporting user querying the view sees something like :

- Product 101 - 5000 units sold, total sales Rs. 12,50,000,500 orders.
- Product 102 - 3200 units sold, total sales Rs. 8,60,000,320 orders.

When the user tries to query the base “sales_transactions” table directly, access is denied, enforcing security.

2. **Objective:** To design and implement secure, efficient data access mechanisms by creating large-scale transaction datasets, summarizing them through normal and materialized views for performance comparison, and enforcing restricted access to sensitive data using views and DCL commands.

3. SQL QUERY AND OUTPUTS -

-----MEDIUM LEVEL PROBLEM-----

Create table TRANSACTION_DATA(id int,val decimal);

INSERT INTO TRANSACTION_DATA(ID,VAL)

SELECT 1,RANDOM()

FROM GENERATE_SERIES(1,1000000);

INSERT INTO TRANSACTION_DATA(ID,VAL)

SELECT 2,RANDOM()

FROM GENERATE_SERIES(1,1000000);

SELECT * FROM TRANSACTION_DATA;



DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

Discover. Learn. Empower.

```
CREATE or REPLACE VIEW SALES_SUMMARY AS  
SELECT  
ID,  
COUNT(*) AS total_quantity_sold,  
sum(val) AS total_sales,  
count(distinct id) AS total_orders  
FROM TRANSACTION_DATA  
GROUP BY ID;
```

```
EXPLAIN ANALYZE  
SELECT * FROM SALES_SUMMARY;
```

```
CREATE MATERIALIZED VIEW SALES_SUMM AS  
SELECT  
ID,  
COUNT(*) AS total_quantity_sold,  
sum(val) AS total_sales,  
count(distinct id) AS total_orders  
FROM TRANSACTION_DATA  
GROUP BY ID;
```

```
EXPLAIN ANALYZE  
SELECT * FROM SALES_SUMM;
```

```
6  
7 INSERT INTO TRANSACTION_DATA(ID,VAL)  
8 SELECT 2,RANDOM()  
9 FROM GENERATE_SERIES(1,1000000);  
10 SELECT * FROM TRANSACTION_DATA;  
11  
12 CREATE or REPLACE VIEW SALES_SUMMARY AS  
13 SELECT  
14 ID,  
15 COUNT(*) AS total_quantity_sold,
```

Data Output Messages Notifications

	id integer	val numeric
1	1	0.0350979238258295
2	1	0.438611872955719
3	1	0.309273924925241
4	1	0.445124939495219
5	1	0.975536437826631



DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

Discover. Learn. Empower.

```
21 EXPLAIN ANALYZE
22 SELECT * FROM SALES_SUMMARY; /*Simple view */
23
```

Data Output					Messages	Notifications
	id	total_quantity_sold	total_sales	total_orders		
	integer	bigint	numeric	bigint		
1	1	1000000	499569.404385168539724401493	1		
2	2	1000000	500134.56553493998172164217	1		

```
20
21 EXPLAIN ANALYZE
22 SELECT * FROM SALES_SUMMARY; /*Simple view */
23
```

Data Output		Messages	Notifications
			Showing rows: 1 to 11
	QUERY PLAN		
	text		
1	GroupAggregate (cost=308493.69..333493.71 rows=2 width=52) (actual time=1458.250..1873.494 rows=2.00 loops=1)		
2	Group Key: transaction_data.id		
3	Buffers: shared hit=10816, temp read=12251 written=12279		
4	-> Sort (cost=308493.69..313493.69 rows=2000000 width=15) (actual time=978.022..1263.665 rows=2000000.00 loops=1)		
5	Sort Key: transaction_data.id		
6	Sort Method: external merge Disk: 49040kB		
7	Buffers: shared hit=10816, temp read=12251 written=12279		
8	-> Seq Scan on transaction_data (cost=0.00..30816.00 rows=2000000 width=15) (actual time=0.282..178.133 rows=2000000.00 loo...		
9	Buffers: shared hit=10816		
10	Planning Time: 0.187 ms		

```
33 EXPLAIN ANALYZE
34 SELECT * FROM SALES_SUMM_MV; /*Materialized view*/
35
```

Data Output					Messages	Notifications
	id	total_quantity_sold	total_sales	total_orders		
	integer	bigint	numeric	bigint		
1	1	1000000	499569.404385168539724401493	1		
2	2	1000000	500134.56553493998172164217	1		

32	
33	EXPLAIN ANALYZE
34	SELECT * FROM SALES_SUMM_MV; /*Materialized view*/
Data Output Messages Notifications	
<div> <div>+</div> <div>📄</div> <div>▼</div> <div>📋</div> <div>▼</div> <div>🗑️</div> <div>🗄️</div> <div>⬇️</div> <div>📈</div> <div>SQL</div> <div>Show</div> </div>	
	QUERY PLAN text 🔒
1	Seq Scan on sales_summ_mv (cost=0.00..20.20 rows=1020 width=52) (actual time=0.022..0.023 rows=2.00 loop...
2	Buffers: shared hit=1
3	Planning Time: 0.099 ms
4	Execution Time: 0.041 ms

OUTPUT -

As we can see that the execution time using the materialized view is very less as compared to the simple view's execution time.

-----HARD PROBLEM-----

```
CREATE TABLE customer_data (
    transaction_id SERIAL PRIMARY KEY,
    customer_name VARCHAR(100),
    email VARCHAR(100),
    phone VARCHAR(15),
    payment_info VARCHAR(50), -- sensitive
    order_value DECIMAL,
    order_date DATE DEFAULT CURRENT_DATE
);
```

-- Insert sample data

```
INSERT INTO customer_data (customer_name, email, phone, payment_info, order_value)
VALUES
('Abhi Kumar', 'abhi@example.com', '9040122324', '1234-5678-9012-3456', 500),
('Sumit Singh', 'sumit@example.com', '9040122324', '1234-5678-9012-3456', 1000),
('Anil Sachdeva', 'anil@example.com', '9876543210', '9876-5432-1098-7654', 700),
('Ashutosh Singh', 'ashutosh@example.com', '9876543210', '9876-5432-1098-7654', 300);
```



DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

Discover. Learn. Empower.

```
CREATE OR REPLACE VIEW RESTRICTED_SALES_DATA AS  
SELECT
```

```
CUSTOMER_NAME,  
COUNT(*) AS total_orders,  
SUM(order_value) as total_sales  
from customer_data  
group by customer_name;
```

```
select * from restricted_sales_data;
```

```
CREATE USER CLIENT1 WITH PASSWORD 'REPORT1234';  
GRANT SELECT ON RESTRICTED_SALES_DATA TO CLIENT1;  
REVOKE SELECT ON RESTRICTED_SALES_DATA FROM CLIENT1;
```

The screenshot shows a SQL IDE interface. At the top, a query is entered in the editor: `select * from restricted_sales_data;`. Below the editor, there are tabs for 'Data Output', 'Messages', and 'Notifications'. The 'Data Output' tab is active, displaying the results of the query in a table format. The table has four columns: 'customer_name' (character varying (100)), 'total_orders' (bigint), and 'total_sales' (numeric). There are four rows of data, numbered 1 to 4 in the first column.

	customer_name character varying (100)	total_orders bigint	total_sales numeric
1	Ashutosh Singh	1	300
2	Abhi Kumar	1	500
3	Sumit Singh	1	1000
4	Anil Sachdeva	1	700



DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

Discover. Learn. Empower.

```
67 GRANT SELECT ON RESTRICTED_SALES_DATA TO CLIENT1;  
68 REVOKE SELECT ON RESTRICTED_SALES_DATA FROM CLIENT1;
```

Data Output Messages Notifications

GRANT

Query returned successfully in 163 msec.

```
64 select * from restricted_sales_data;  
65  
66 CREATE USER CLIENT1 WITH PASSWORD 'REPORT1234';  
67 GRANT SELECT ON RESTRICTED_SALES_DATA TO CLIENT1;  
68 REVOKE SELECT ON RESTRICTED_SALES_DATA FROM CLIENT1;
```

Data Output Messages Notifications

SQL

	customer_name character varying (100) 🔒	total_orders bigint 🔒	total_sales numeric 🔒
1	Ashutosh Singh	1	300
2	Abhi Kumar	1	500
3	Sumit Singh	1	1000
4	Anil Sachdeva	1	700



DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

Discover. Learn. Empower.

```
66 CREATE USER CLIENT1 WITH PASSWORD 'REPORT1234';
67 GRANT SELECT ON RESTRICTED_SALES_DATA TO CLIENT1;
68 REVOKE SELECT ON RESTRICTED_SALES_DATA FROM CLIENT1;
```

Data Output Messages Notifications

REVOKE

Query returned successfully in 152 msec.

```
80
81 -- Select from restricted view
82 SELECT * FROM RESTRICTED_SALES_DATA;
83
84 -- Create user and grant/revoke access
```

Data Output Messages Notifications

ERROR: permission denied for table customer_data

SQL state: 42501